

TTIC 31230 Fundamentals of Deep Learning

Problems for Graphical Models.

Problem 1. Dynamic Programing for HMMs Assume we have an input sequence x_1, \dots, x_T and a phoneme gold label y_1, \dots, y_T with $y_t \in \mathcal{P}$. This problem is simpler than CTC because the gold label has the same length as the input sequence.

In an HMM we assume a hidden state sequence s_1, \dots, s_T with $s_t \in \mathcal{S}$ where \mathcal{S} is some finite sets of “hidden states”. Here will assume that then some deep network has computed transition probabilities and emission probabilities.

$$P_{\text{Trans}}(s_{t+1} \mid s_t)$$

$$P_{\text{Emit}}(y_t \mid s_t)$$

We assume an initial state s_{init} and a stop state s_{stop} such that $s_1 = s_{\text{init}}$ (before emitting any phonemes). The length T is determined by when the hidden state becomes s_{stop} giving $s_{T+1} = s_{\text{stop}}$.

For a given gold sequence y_1, \dots, y_T we define a “forward tensor” as

$$F[t, s] = P(y_1, \dots, y_{t-1} \wedge s_t = s)$$

We have

$$\begin{aligned} F[1, s_{\text{init}}] &= 1 \\ F[1, s] &= 0 \quad \text{for } s \neq s_{\text{init}} \end{aligned}$$

(a) Write a dynamic programming equation to compute $F[t, s]$ from $F[t-1, s']$ for various values of s' .

Solution:

$$F[t, s] = \sum_{s'} F[t-1, s'] P_{\text{Emit}}(y_{t-1} \mid s') P_{\text{Trans}}(s \mid s')$$

(b) Express $P(y_1, \dots, y_T)$ in terms of $F[t, s]$.

Solution:

$$P(y_1, \dots, y_T) = F[T+1, s_{\text{stop}}]$$

(c) EM for HMMs involves computing a “backward” tensor

$$B[t, s] = P(y_t, \dots, y_T \mid s_t = s).$$

Explain why, if the forward equations are written in a framework, we do not need to also compute the backward tensor.

Solution: Once we have expressed the loss $-\ln P(y_1, \dots, y_T)$ in a framework we can train the model by SGD using the framework’s implementation of back-propagation. Nothing more is needed.

Problem 2. CTC for image labeling

Suppose that the training data consists of pairs (I, S) where I is an image and S is a set of object types occurring in the image. For example S might be $\{\text{Person, Dog, Car}\}$. To be concrete we can take \mathcal{C} to be the set of image labels used in CIFAR 100 and take S to be a subset of \mathcal{C} containing no more than five labels ($|S| \leq 5$). We want to do SGD on a model defining $P_\Phi(S \mid I)$.

We will use a latent variable $z[X, Y]$ such that for pixel coordinates (x, y) we have $z[x, y] \in \mathcal{C} \cup \{\perp\}$. For a given $z[X, Y]$ define $S(z[X, Y])$ to be the set of classes appearing in $z[X, Y]$, i.e., $S(z[X, Y]) = \{c \mid \exists x, y \ z(x, y) = c\}$. Here the “semantic segmentation” $Z[X, Y]$ is analogous to the phoneme sequence $z[T]$ in CTC. Unlike the CTC model, the label S is a set rather than a sequence.

We assume a CNN (with convolutions of stride 1 to preserve spatial dimensions) followed by a softmax at each pixel to get a probability $P_\Phi(z[x, y] = c)$ for each pixel location (x, y) and each $c \in \mathcal{C} \cup \{\perp\}$ and where each pixel location has an independent probability distribution over classes. To simplify notation we can reshape the pixel locations into a linear sequence and replace $z[X, Y]$ by $z[T]$ with $T = X \times Y$ so we have $z[1], z[1], \dots, z[T]$.

Define

$$S_t = \{c \in \mathcal{C} \mid \exists t' \leq t \ z[t'] = c\}$$

For $U \subseteq S$ define

$$F[U, t] = P(S_t = U)$$

Note that for $|S| \leq 5$ there are at most 32 possible values of U . Give dynamic programming equations defining $F[U, 0]$ and defining $F[U, t + 1]$ in terms of $F[U', t]$ for various U' .

Solution:

$$F[\emptyset, 0] = 1$$

$$\text{For } U \text{ a nonempty subset of } S \ F[U, 0] = 0$$

$$\text{For } t = 1, \dots, T$$

$$\text{For } U \subseteq S$$

$$F[U, t] = P(z[t] = \perp)F[U, t - 1] + \sum_{c \in U} P(z[t] = c)(F[U \setminus c, t - 1] + F[U, t - 1])$$

Problem 3. Pseudolikelihood of a one dimensional spin glass. We let \hat{x} be an assignment of a value to every node where the nodes are numbered from 1 to N_{nodes} and for every node i we have $\hat{x}[i] \in \{0, 1\}$. We define the score of \hat{x} by

$$f(\hat{x}) = \sum_{i=1}^{N-1} \mathbf{1}[\hat{x}[i] = \hat{x}[i+1]]$$

The probability distribution over assignments is defined by a softmax. We let $\hat{x}[i := v]$ be the assignment identical to \hat{x} except that node i is assigned the value v . The expression $\hat{x}[i] = v$ is either true or false depending on whether node i is assigned value v in \hat{x} . So these are quite different.

$$P_f(\hat{x}) = \underset{\hat{x}}{\text{softmax}} f(\hat{x})$$

Pseudolikelihood is defined in terms of the softmax probability P_f as follows.

$$\tilde{P}_f(\hat{x}) = \prod_i P_f(\hat{x}[i] \mid \hat{x} \setminus i)$$

What is the pseudolikelihood of the all ones assignment under the definition of f given above?

Solution: In a graphical model $P_f(\hat{x}[i] \mid \hat{x}/i)$ is determined by the neighbors of i and we can consider only how a value is scored against it neighbors. For \hat{x} equal to all ones we have

$$f(\hat{x}) = N - 1$$

$$f(\hat{x}[i := 0]) = \begin{cases} N - 3 & \text{for } 1 < i < N \\ N - 2 & \text{for } i = 1 \text{ or } i = N \end{cases}$$

For $1 < i < N$ we get

$$\begin{aligned} Q_f(\hat{x}[i = 1] \mid \hat{x}/i) &= \frac{e^{N-1}}{e^{N-1} + e^{N-3}} \\ &= \frac{1}{1 + e^{-2}} \end{aligned}$$

and for $i = 1$ or $i = N$ we get

$$Q_f(\hat{x}[i = 1] \mid \hat{x}/i) = \frac{1}{1 + e^{-1}}$$

This gives

$$\tilde{Q}(\hat{x}) = (1 + e^{-1})^{-2} (1 + e^{-2})^{-(N-2)}$$

Problem 4. Pseudolikelihood for images. Consider a semantic segmentation $\hat{y}[i]$ on pixels i with $\hat{y}[i]$ a semantic class label in $\{C_1, \dots, C_K\}$. We also assume a scoring function s_Φ on semantic segmentations defining

$$P_\Phi(\hat{y}) = \underset{\hat{y}}{\text{softmax}} \ s_\Phi(\hat{y})$$

Pseudolikelihood is defined by

$$\tilde{P}_\Phi(\hat{y}) = \prod_i P_\Phi(\hat{y}[i] \mid \hat{y} \setminus i)$$

where $\hat{y} \setminus i$ assigns a class to every pixel other than i , and $\hat{y}[i := c]$ is the semantic segmentation identical to \hat{y} except that pixel i is labeled with semantic class c . In a typical graphical model for images we have

$$P_\Phi(\hat{y}[i] \mid \hat{y} \setminus i) = P_\Phi(\hat{y}[i] \mid \hat{y}[N(i)])$$

where $\hat{y}[N(i)]$ is \hat{y} restricted to those pixels which are neighbors of pixel i .

(a) show

$$\frac{P_\Phi(\hat{y})}{\sum_c P_\Phi(\hat{y}[i := c])} = \underset{c}{\text{softmax}} \ s_\Phi(\hat{y}[i := c]) \quad \text{evaluated at } c = y[i]$$

Solution:

$$\begin{aligned} \frac{P_\Phi(\hat{y})}{\sum_c P_\Phi(\hat{y}[i := c])} &= \frac{\frac{1}{Z} e^{s_\Phi(\hat{y})}}{\sum_c \frac{1}{Z} e^{s_\Phi(\hat{y}[i := c])}} \\ &= \frac{e^{s_\Phi(\hat{y})}}{\sum_c e^{s_\Phi(\hat{y}[i := c])}} \\ &= \underset{c}{\text{softmax}} \ s_\Phi(\hat{y}[i := c]) \quad \text{evaluated at } c = y[i] \end{aligned}$$

(b) How many scores need to be computed in the worst case for computing $P_\Phi(\hat{y})$. Given the result of part (a), how many for computing $\tilde{P}_\Phi(\hat{y})$?

Solution: K^N for P_Φ and KN for \tilde{P}_Φ .

(c) Consider a distribution on semantic segmentations where for each pixel the class assigned to that pixel is uniquely determined by the classes of its neighbors. Can this distribution be defined by a softmax over scores? Explain your answer.

Solution: No. Since $e^s > 0$ for any finite s , all semantic segmentations must have nonzero probability.

(d) If P_Φ is a distribution defined in some other way such that the class of each pixel is completely determined by the other pixels, given a simple expression for $\tilde{P}_\Phi(\hat{y})$ in the case where \hat{y} has nonzero probability under P_Φ .

Solution: We have $P_\Phi(\hat{y}|\hat{y}\setminus i) = 1$ which implies $\tilde{P}(\hat{y}) = 1$.

Problem 5. Pseudolikelihood for Monocular Distance Estimation.

(25 points) Here we are interested in labeling each pixel with a distance from the camera. Each pixel i is to be labeled with a real number $\hat{y}[i] > 0$ giving the distance in (say) meters from the camera to the point on the object displayed by that pixel. We assume a neural network that computes for each pixel i an expected distance μ_i and a variance $\sigma_i > 0$. For each pair of neighboring pixels i and j the neural network computes a real number $\lambda_{\langle i, j \rangle} \geq 0$. For each assignment \hat{y} of distances to pixels we then define the score $s(\hat{y})$ by

$$s(\hat{y}) = \sum_{i \in \text{nodes}} -(\hat{y}[i] - \mu_i)^2 / \sigma_i^2 + \sum_{\langle i, j \rangle \in \text{edges}} -\lambda_{\langle i, j \rangle} |\hat{y}[i] - \hat{y}[j]|$$

(a) This scoring function determines a continuous softmax distribution defined by

$$p(\hat{y}) = \frac{1}{Z} e^{s(\hat{y})}$$

where Z is an integral rather than a sum. What is the dimension of the space to be integrated over in computing Z ?

Solution: This is an integration over \mathbb{R}^N where N is the number of nodes — an N_{nodes} dimensional space.

(b) We now consider pseudolikelihood for this problem. Give an expression for the continuous conditional probability density on $\hat{y}[i]$ for the distance $\hat{y}[i]$ conditioned on the value of the neighbors $N(i)$ of node i . This probability is written $p(\hat{y}[i] | \hat{y}[N(i)])$. Your answer should be given as a function of the values $\hat{y}[j]$ for the nodes j neighboring i written $j \in N(i)$. Write Z as an integral but do not bother trying to solve it. What is the dimension of the integral for this conditional probability?

Solution:

$$\begin{aligned} p(\hat{y}[i] \mid \hat{y}[N(i)]) &= \frac{1}{Z} \exp \left(-(\hat{y}[i] - \mu_i)^2 / \sigma_i^2 + \sum_{j \in N(i)} -\lambda_{\langle i, j \rangle} |\hat{y}[i] - \hat{y}[j]| \right) \\ Z &= \int_0^\infty \exp \left(-(x - \mu_i)^2 / \sigma_i^2 + \sum_{j \in N(i)} -\lambda_{\langle i, j \rangle} |x - \hat{y}[j]| \right) dx \end{aligned}$$

This is an integral over a one dimensional space (a single real number).